University Student-Support Case Agent

#	Capability	AI May Do (Autonomous)	Stays Deterministic (Rule-Based Code)	Requires Human Approval
1	Answering student questions	Interpret the question, retrieve relevant passages from approved handbook/course documents, and generate a grounded answer with citation to the source document.	The retrieval/search logic itself (which documents are indexed, similarity thresholds, ranking) is fixed code, not model judgment.	None required for a grounded answer. If the agent cannot find a grounded answer, it must state this rather than guess fabricated answers.
2	Checking timetable / case status	Interpret the student's request and call the lookup tool, then phrase the result in plain language for the student.	The actual lookup (querying the timetable/case-status data) is a deterministic function call, not an LLM guess.	None; this is a read-only, low-risk action.
3	Creating / routing a support ticket	Draft the ticket content (summary, category, priority) from the conversation.	Ticket ID generation, routing rules (which department/queue), and storage are deterministic.	The student confirms before the ticket is actually submitted, since this creates a record and starts a process.
4	Remembering the active case	Decide what is relevant to carry forward in the current session (e.g. which	What fields are stored, retention period, and access permissions are	None for normal use, but stored case memory should be periodically audited by

#	Capability	AI May Do (Autonomous)	Stays Deterministic (Rule-Based Code)	Requires Human Approval
		case is being discussed).	fixed by design, not by the model.	a human for data that should not be retained.
5	Admissions, grading, disciplinary, or fee decisions	Never. The agent refuses and redirects the student to the correct human office.	These decision types are excluded from the agent's scope entirely, by design.	Always. Any request touching these areas is deflected to a human/office and no AI judgment is applied.
6	Escalating unclear or distressed cases	May flag or recommend escalation based on conversation cues.	The escalation trigger rule (e.g. specific keywords, repeated failed lookups) is deterministic, not model-invented.	A human staff member reviews and decides on the escalation; the agent does not escalate unilaterally.