

University Student-Support Case Agent

AI Boundary Matrix

#	Capability	AI May Do (Autonomous)	Stays Deterministic (Rule-Based Code)	Requires Human Approval
1	Answering student questions	Interprets the question, decides when document search is required, and generates a grounded answer from the passages returned by the search tool.	The retrieval/search logic itself (which documents are indexed, similarity thresholds, ranking) is fixed code, not model judgment.	None required for a grounded answer. If the agent cannot find a grounded answer, it must state this rather than guess; no fabricated answers are permitted.
2	Checking timetable / case status	Interpret the student's request and call the lookup tool, then phrase the result in plain language for the student.	The actual lookup (querying the timetable/case-status data) is a deterministic function call, not an LLM guess.	None; this is a read-only, low-risk action.
3	Creating / routing a support ticket	Draft the ticket content (summary, category, priority) from the conversation.	Ticket ID generation, routing rules (which department/queue), and storage are deterministic.	The student confirms before the ticket is actually submitted, since this creates a record and starts a process.
4	Remembering the active case	Decide what is relevant to carry forward in the current session (e.g. which case is being discussed).	What fields are stored, retention period, and access permissions are fixed by design, not by the model.	Student consent is required before information is retained beyond the active session.
5	Admissions, grading, disciplinary, or fee decisions	Never. The agent refuses and redirects the student to the correct human office.	These decision types are excluded from the agent's scope entirely, by design.	Human handling is required. The agent cannot make or recommend these decisions.
6	Escalating unclear or sensitive cases	May flag or recommend escalation based on conversation cues.	The escalation trigger rule (e.g. specific keywords, repeated failed lookups) is deterministic, not model-invented.	A human staff member reviews and decides on the escalation; the agent does not escalate unilaterally.